

[Return to "Machine Learning Engineer Nanodegree" in the classroom](#)

Teach a Quadcopter How to Fly

REVIEW

CODE REVIEW

HISTORY

Meets Specifications

- I would like to highly recommend you to go through the following links for better clarity on DDPG and how to tackle similar RL problems.
 - <http://pemami4911.github.io/blog/2016/08/21/ddpg-rl.html>
 - <https://towardsdatascience.com/reinforcement-learning-w-keras-openai-actor-critic-models-f084612cfd69>
 - <https://medium.com/@BonsaiAI/deep-reinforcement-learning-models-tips-tricks-for-writing-reward-functions-a84fe525e8e0>

Define the Task, Define the Agent, and Train Your Agent!

The `agent.py` file contains a functional implementation of a reinforcement learning algorithm.

Awesome

- Effectively actualised both the Actor and Critic network for DDPG.
- The implementation is highly modular, making it simple to troubleshoot and effortlessly scalable.

The `Quadcopter_Project.ipynb` notebook includes code to train the agent.

Awesome

- The code to train the agent is attached.

Plot the Rewards

A plot of rewards per episode is used to illustrate how the agent learns over time.

Awesome

- The plot of rewards is attached.

Reflections

The submission describes the task and reward function, and the description lines up with the implementation in `task.py`. It is clear how the reward function can be used to guide the agent to accomplish the task.

Awesome

- Great job with the task.
- The whole task was quite well written with all the different functions and conditions defined clearly.
- The reward function was also quite well explained in the Jupyter notebook.

The submission provides a detailed description of the agent in `agent.py`.

Awesome

- The choice of parameters was good. The choice of using DDPG for the agent's continuous action space was a good choice.
- It was good to see that you didn't just use the default set of hyperparameters.
- Manipulation of hyperparameters is a great learning exercise and helps to understand its impact on the performance of the agent.

The submission discusses the rewards plot. Ideally, the plot shows that the agent has learned (with episode rewards that are gradually increasing). If not, the submission describes in detail various attempted settings (hyperparameters and architectures, etc) that were tested to teach the agent.

- As per the rubric requirements, the agent doesn't have to absolutely learn to complete the project successfully.
- It seems to be the case in your submission, but you have provided detail regarding the various attempted settings (hyperparameters and architectures, etc) that were tested to teach the agent.
- It is sufficient to complete this rubric point successfully.

A brief overall summary of the experience working on the project is provided, with ideas for further improving the project.

Awesome

- You did a great job on this project.
- The code was quite modular which made it easy to read and debug.
- You have a keen eye on the domain knowledge and required parameters to get the job done.
- All the points in the Jupyter notebook were quite well articulated.

 [DOWNLOAD PROJECT](#)

[RETURN TO PATH](#)
